ANOOP REDDY YEDDULA

Portland, OR | mailto:anoopreddy3001@gmail.com | (917)-624-1889 | LinkedIn | GitHub | Hugging Face

EDUCATION

Master of Science, Computer Science September 2022 - June 2024
Portland State University, Portland, Oregon
Bachelor of Engineering, Computer Science August 2016 - May 2020
Dayananda Sagar University, Bengaluru, Karnataka, India

TECHNICAL SKILLS

Programming & Web Frameworks: Python, Java, C/C++, JavaScript, HTML/CSS, React.js, Node.js, FastAPI, Bootstrap
AI/ML & Data: TensorFlow, PyTorch, Vertex AI, Scikit-learn, LangChain, BigQuery, Data Pipelines, RAG workflows
Cloud & DevOps: GCP, AWS, Azure, Docker, Kubernetes, Terraform, Azure DevOps
Tools & IDEs: Git, GitHub, Jupyter, Eclipse, MATLAB, Visual Studio, Android Studio
APIs & Visualization: REST APIs, API Gateway, Power BI, Tableau, Model Monitoring & Retraining
Certifications (Links): AWS CERTIFIED CLOUD PRACTITIONER, Build and Deploy Machine Learning Solutions on Vertex AI

PROFESSIONAL EXPERIENCE

AI/ML Engineer | KKRGenAI Innovations LLC, United States October 2024 - Present

- Designed and deployed scalable **ML models** using **Vertex AI**, **TensorFlow**, and **PyTorch**, achieving **96% accuracy** in finance and healthcare.
- Built and maintained **BigQuery** data pipelines, reducing **ETL processing time by 40%** and enabling **real-time analytics**.
- Delivered real-time predictions through **FastAPI** and **API Gateway**, supporting **10K+ daily requests** with minimal latency.
- Containerized microservices with **Docker** and orchestrated with **Kubernetes** on **GCP**, achieving **99.9% uptime**.
- Integrated **RAG pipelines** using **LangChain**, boosting **LLM response relevance by 35%**.
- Implemented **CI/CD workflows** with **Cloud Build** and **GitHub Actions**, reducing deployment time by **50%**.

Data Analyst | Hudl, India October 2020 - May 2022

- Analyzed player and team performance across **500+ matches**, boosting win rates by **15%**.
- Automated reporting with **Python** (**Pandas, NumPy, Matplotlib**), reducing manual tasks by **60%** and improving reporting frequency.
- Developed advanced **SQL queries** (joins, CTEs, window functions) to extract patterns from large datasets, enabling **real-time insights**.
- Created interactive dashboards with **Power BI** and **Tableau**, used by **50+ coaches and analysts**.
- Collaborated with engineers to embed analytics into Hudl's platform, expanding user access to data insights.
- Contributed to AI model integration, increasing **scouting efficiency by 30%** and maintaining **98%+ data accuracy**.

PROJECTS (GitHub)

Customer Churn Prediction using LLM and ML | *LangChain, Python, Hugging Face, Scikit-learn* (GitHub)

- Leveraged **scikit-learn** and **Hugging Face**'s **flan-t5-base LLM** to predict customer churn using the **BankChurners dataset**, enhancing model accuracy by 30%.
- Integrated **LangChain** with **LLM workflows** to generate human-like insights for customer retention strategies.
- Successfully deployed an end-to-end solution that predicts churn and provides actionable insights for businesses.

Automated Insurance Claim Validation System | *NLP, Python, Image Processing, Gradio* (GitHub)

- Automated insurance claim validation using **EasyOCR** for text extraction and **BERT** for document classification, improving validation accuracy by 35% and reducing manual review time by 40%.
- Integrated a multi-page processing system that supports both **PDFs** and **image files**, streamlining claim processing and enhancing operational efficiency. Deployed as a **Gradio** ensuring seamless end-user interactions with a scalable backend.

Real-Time Inventory Advisor | *GenAI, ML, Gradio, Python* (GitHub)

- Developed an **AI-powered inventory management system** using **Gradio**, **Scikit-learn**, and **Hugging Face Transformers**, predicting sales trends and providing real-time restocking recommendations.
- Achieved 95% accuracy in restocking predictions, reducing stockouts by 25% and optimizing product availability.
- Integrated **predictive modeling** with **generative AI** for dynamic inventory forecasting and decision-making.

Credit Score Type Prediction | *ML, Python, Streamlit* (GitHub)

- Developed a **machine learning application** to predict **credit score types** (Poor, Average, Good) based on **financial indicators**, utilizing real-time data for accurate predictions.
- Utilized **Streamlit** to build an interactive web application, enabling users to input financial data and receive instant credit score predictions.
- Implemented a robust model using **Random Forest**, achieving high predictive accuracy in financial assessments.

Drug Effectiveness Prediction | *ML, Python, Logistic Regression, TF-IDF, Gradio* (GitHub)

- Developed a system to predict drug effectiveness for various medical conditions using **logistic regression** and **TF-IDF vectorization**, achieving an 88% classification accuracy.
- Streamlined medical decision-making by providing actionable insights into drug performance based on real-world data.
- Integrated a user-friendly **Gradio interface** to allow easy input of drug names and medical conditions for instant predictions.

Decoding Facial Recognition using CNN | *Python, OpenCV, TensorFlow* (GitHub)

- Led **CNN architecture design (40%)** and contributed to **data preprocessing (30%)**, building a model with **95% facial recognition accuracy** on custom datasets.
- Optimized performance using **TensorFlow** and **Keras**; visualized training metrics to fine-tune results and improve **model evaluation (30%)**.